

Final Project - Travel Insurance Predictions Team 7

October 8, 2021

ADS-505 Final Project - Travel Insurance Predictions

Team: #7

Team Members: Jimmy Nguyen, Christopher Robinson, Nima Amin Taghavi

Date: 09/20/2021

Programmin Language: Python Code

Table of Contents

- 1 Problem statement
- 2 Packages
- 3 Data Set
- 4 Exploratory Data Analysis (EDA)
 - 4.1 Exploring Response Variable - TravelInsurance
 - 4.2 Examining customers' Age
 - 4.2.1 Age Distribution
 - 4.2.2 Age with Target Variable Overlaid
 - 4.2.3 Normalized Histogram with Target Variable Overlaid on Age
 - 4.2.4 Age Groups Comparison (20s vs. 30s)
 - 4.2.5 Percentage of Travel Insurance Purchases on Various Features
 - 4.3 Monthly Income (Indian Rupee) Range
 - 4.4 Side-by-side Box-plots between Annual Income and Different Attributes
- 5 Feature Engineering and Pre-Processing
- 6 Data splitting
- 7 Model building strategies
- 8 Model performance and hyper-parameter tuning
 - 8.1 Decision Tree

8.2	Logistic Regression
8.3	Multi-Layered Neural Network
9	Results and final model selection
10	Discussion and conclusion

1 Problem statement

About the Client

The client in this data mining project is a tour & travels company that is offering travel insurance package to their customers. This new insurance package also includes COVID-19 coverage for their flights. However, the client wants to know which customers based on their data base history are potential purchasers who may be interested in buying this new insurance package. Previously, the insurance package was offered to some of the customers in 2019 and data was collected from the performance and sales of the package during that period. The sample data given has close to 2000 customers from that period. The client is requesting information on which customer are most likely going to buy the travel insurance given their information such as employment type, income level, etc.

Business Problem

The client may use the solutions presented to them for customer-targeted advertising of the new travel insurance package. Also, data visualizations provided will help derive interesting insights about their potential buyers in order to optimize marketing strategies.

Data Mining Problem

- A supervised classification task, where the outcome variable of interest is *TravelInsurance* that indicates whether the customer will buy the travel insurance. Performance metrics should take in consideration the positive class of buyers/purchasers.
 - Find out interesting patterns and trends for better customer segmentations through data exploration and visualizations.
 - An unsupervised task, where the goal is to cluster customers.
-

2 Packages

Python code:

```
[1]: %%javascript
IPython.OutputArea.prototype._should_scroll = function(lines) {
    return false;
}
```

<IPython.core.display.Javascript object>

```
[2]: from pathlib import Path
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.linear_model import LinearRegression, LogisticRegression,
↳ LogisticRegressionCV
from sklearn.model_selection import train_test_split
import statsmodels.api as sm
import scikitplot as skplt
from mord import LogisticIT
from sklearn import preprocessing
from sklearn.metrics import accuracy_score
from sklearn.neighbors import KNeighborsClassifier
from sklearn.tree import DecisionTreeRegressor, DecisionTreeClassifier
from sklearn.neural_network import MLPClassifier
from sklearn.model_selection import cross_val_score, GridSearchCV
from dmba import regressionSummary, stepwise_selection, plotDecisionTree
from dmba import regressionSummary, stepwise_selection
from dmba import forward_selection, backward_elimination, exhaustive_search
from dmba import classificationSummary, gainsChart, liftChart
from dmba.metric import AIC_score
from tabulate import tabulate
import matplotlib.patches as mpatches
import warnings
sns.set_theme()
plt.rcParams['figure.figsize'] = [11, 9]

warnings.filterwarnings('ignore')
```

3 Data Set

Data Dictionary

1. **Age** - Age Of The Customer
2. **Employment Type** - The Sector In Which Customer Is Employed
3. **GraduateOrNot** - Whether The Customer Is College Graduate Or Not
4. **AnnualIncome** - The Yearly Income Of The Customer In Indian Rupees[Rounded To Nearest 50 Thousand Rupees]
5. **FamilyMembers** - Number Of Members In Customer's Family
6. **ChronicDisease** - Whether The Customer Suffers From Any Major Disease Or Conditions Like Diabetes/High BP or Asthama,etc.

7. **FrequentFlyer** - Derived Data Based On Customer's History Of Booking Air Tickets On Atleast 4 Different Instances In The Last 2 Years[2017-2019].
8. **EverTravelledAbroad** - Has The Customer Ever Travelled To A Foreign Country[Not Necessarily Using The Company's Services]
9. **TravelInsurance** - Did The Customer Buy Travel Insurance Package During Introductory Offering Held In The Year 2019.

Python code:

```
[3]: # Load data set
df = pd.read_csv("../Data/TravelInsurancePrediction.csv")

# First few rows of data set
df.head(3)
```

```
[3]:   Age      Employment Type GraduateOrNot  AnnualIncome \
0   31      Government Sector             Yes      400000
1   31  Private Sector/Self Employed       Yes     1250000
2   34  Private Sector/Self Employed       Yes      500000

      FamilyMembers  ChronicDiseases FrequentFlyer EverTravelledAbroad \
0                6                1           No                No
1                7                0           No                No
2                4                1           No                No

      TravelInsurance
0                0
1                0
2                1
```

4 Exploratory Data Analysis (EDA)

- Graphical and non-graphical representations of relationships between the response variable and predictor variables

4.1 Exploring Response Variable - TravelInsurance

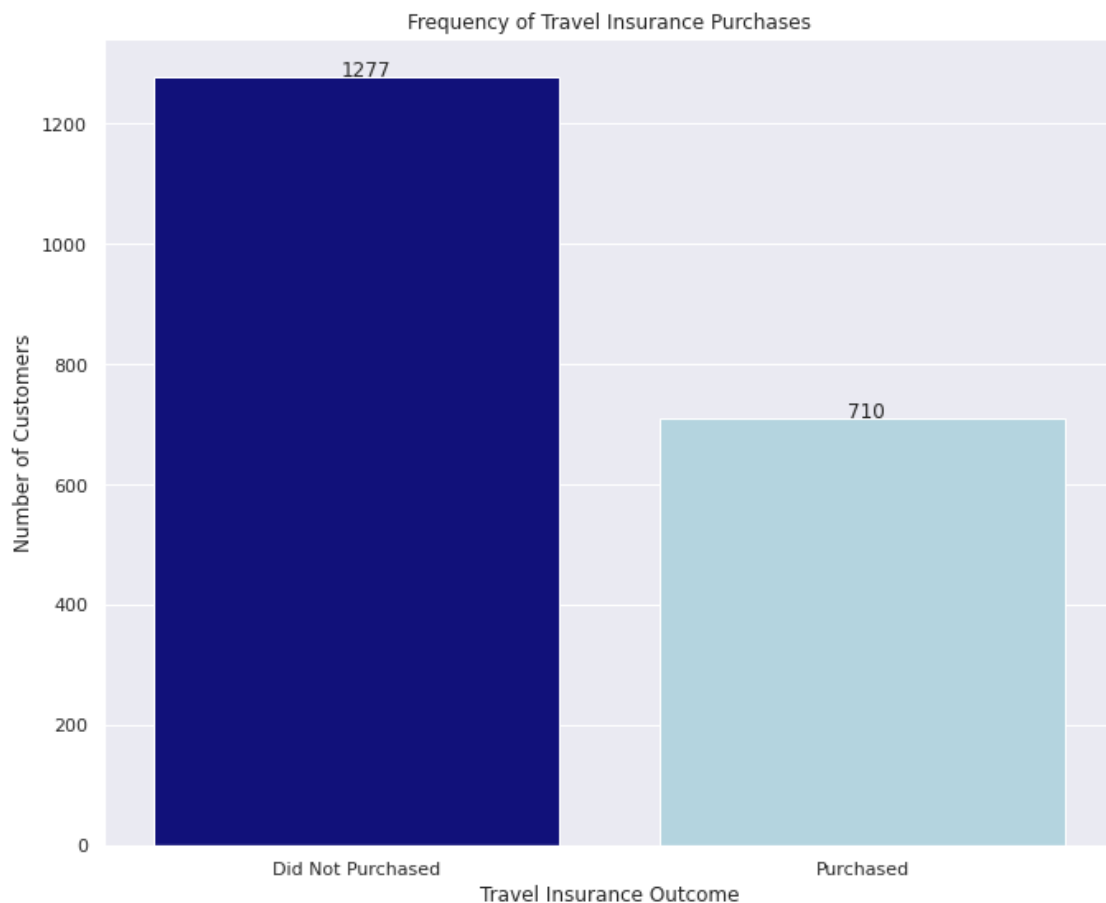
- Examine the frequency of travel insurance purchases
- 0 = No
- 1 = Yes

Python code:

```
[4]: # Graph count plot of Non-Purchases vs. Purchases
ax = sns.countplot(data=df, x="TravelInsurance", palette = ["darkblue", "lightblue"])

# Add count labels
for p, label in zip(ax.patches, df['TravelInsurance'].value_counts()):
    ax.annotate(label, (p.get_x()+0.37, p.get_height()+0.15))

# Graph Properties
plt.title("Frequency of Travel Insurance Purchases")
plt.xticks([0, 1], ["Did Not Purchased", "Purchased"])
plt.xlabel("Travel Insurance Outcome")
plt.ylabel('Number of Customers')
plt.show()
```



Summary

- There is a higher number of customers who flew without buying travel insurance in this sample data set.

4.2 Examining customers' Age

- Age distributions
- Age with Target Variable Overlaid
- Normalized Histogram with Target Overlaid on Age
- Age Group Comparisons (20s vs. 30s)
- Percentage of Purchases between Age groups (20s vs. 30s)

4.2.1 Age Distribution

Python code:


```
[5]: # Get a range of customer ages
age = pd.DataFrame({'Age': df['Age'].value_counts().sort_index()})
age
```

```
[5]:      Age
25    146
26    148
27    131
28    506
29    192
30     65
31    234
32     91
33    148
34    266
35     60
```

```
[6]: # Histogram of Age and set the range of bins from 25-35
bins = np.arange(25, 36)
ax = df['Age'].plot.hist(bins=bins)

# add labels
for p, label in zip(ax.patches, df['Age'].value_counts().sort_index()):
    ax.annotate(label, (p.get_x() + 0.37, p.get_height() + 0.15))

# title and axis
plt.title("Histogram of Customer Ages")
plt.xlabel("Age Range of Customers")
plt.ylabel("Number of Customers")
plt.show()
```



Summary

- There are 506 customers who are 30 years-old which is visualized as the most in this data set
- While customers who are 36 years-old are the least in this data set.

4.2.2 Age with Target Variable Overlaid

Python code:

```
[7]: # Set up plot with response overlaid
n, bins, patches = plt.hist(
    [
        df[df['TravelInsurance'] == 1]['Age'],
        df[df['TravelInsurance'] == 0]['Age']
    ],
    bins=10,
    stacked=True,
    color=["lightblue", "darkblue"],
)
```

```
# title and axis
labels = ["Travel Insurance Purchased", "Non-purchase"]
plt.legend(labels)
plt.title("Histogram of Customer Ages Travel Insurance Overlaid")
plt.xlabel("Age Range of Customers")
plt.ylabel('Number of Customers')
plt.show()
```



Summary

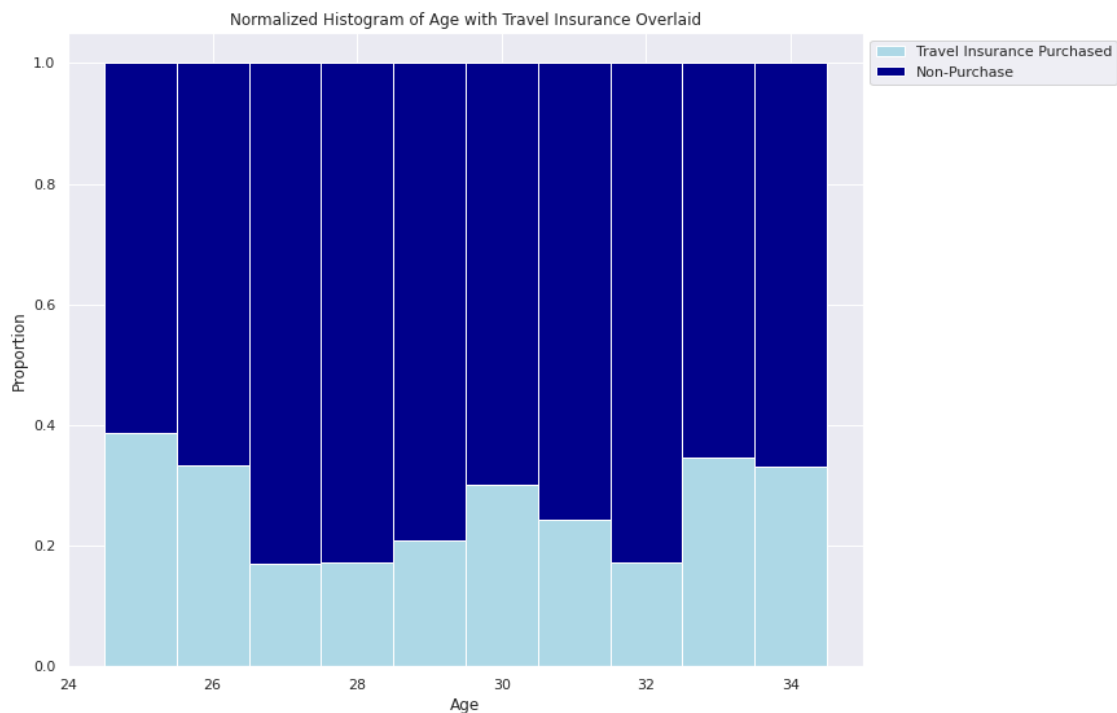
- It is difficult to compare between age groups with target variable overlaid
- Therefore, it is better to focus on one class from the target variable and analyze age in a normalized histogram.
- The following is visualized below.

4.2.3 Normalized Histogram with Target Variable Overlaid on Age

Python code:


```
[8]: # Create normalized histogram for age groups by target overlay
n_table = np.column_stack((n[0], n[1])) # stack the tables
n_norm = n_table / n_table.sum(axis=1)[:,
                                     None] # create normalized tables by sum
ourbins = np.column_stack((bins[0:10], bins[1:11])) # create table bins

p1 = plt.bar(x=ourbins[:, 0],
             height=n_norm[:, 0],
             width=ourbins[:, 1] - ourbins[:, 0], color = "lightblue") # first
→ bar chart
p2 = plt.bar(
    x=ourbins[:, 0],
    height=n_norm[:, 1],
    width=ourbins[:, 1] - ourbins[:, 0], # second bar chart
    bottom=n_norm[:, 0], color = "darkblue")
# Annotate legend, title with x and y labels
plt.legend(['Travel Insurance Purchased', 'Non-Purchase'],
           bbox_to_anchor=(1, 1))
plt.title('Normalized Histogram of Age with Travel Insurance Overlaid')
plt.xlabel('Age')
plt.ylabel('Proportion')
plt.show()
```



Summary

- Insights for this graph show that it may be better to compare age classified into 2 groups instead such as customers who are in their twenties (20s) vs customers in their thirties (30s).

4.2.4 Age Groups Comparison (20s vs. 30s)

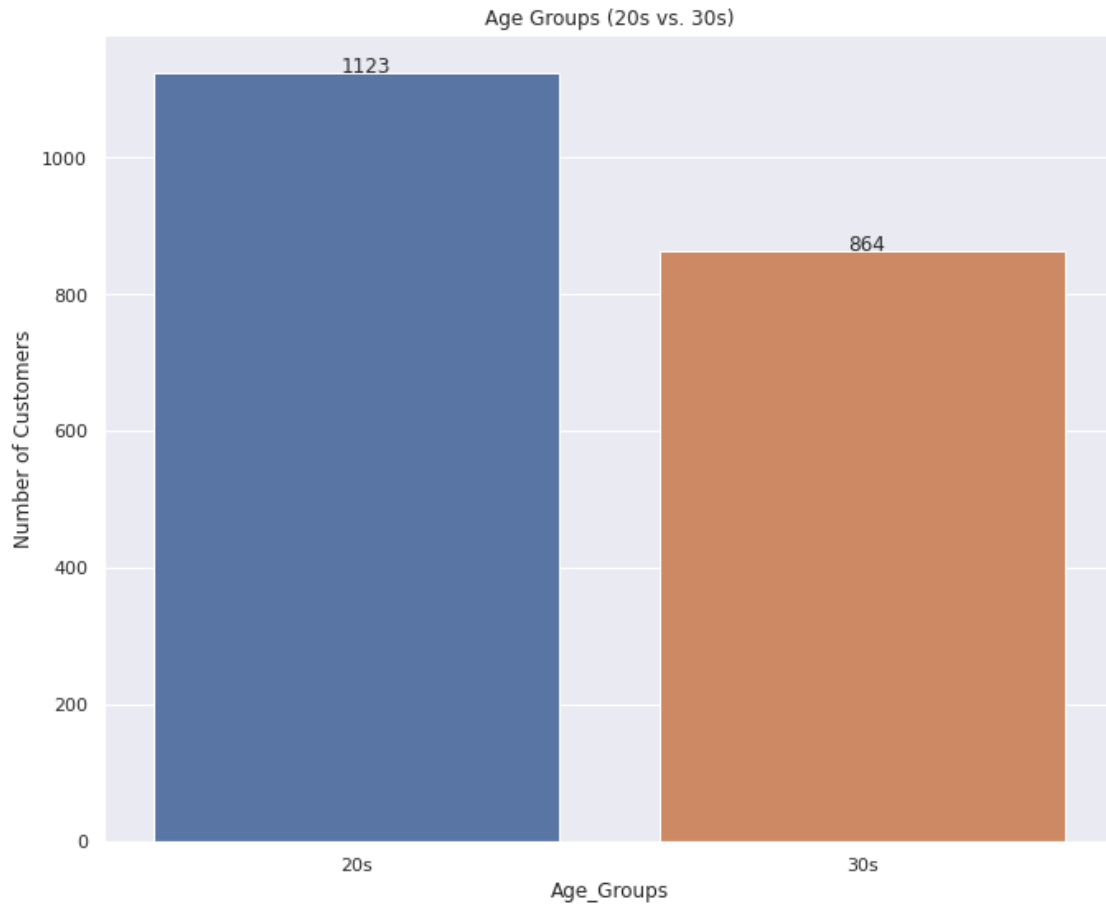
Python code:

```
[9]: # Create function to categorize age groups
def age_groups(x):
    """
    x: This is a value from df['Age']
    returns each as a new categorical value of 20s or 30s
    """
    if x < 30:
        return '20s'
    else:
        return '30s'

# Apply age_groups function on each value
age_groups = pd.DataFrame(
    {'Age_Groups': df['Age'].apply(lambda x: age_groups(x)),
     'AnnualIncome':df['AnnualIncome'],
     'TravelInsurance':df['TravelInsurance']})

# Graph count plot of age groups (20s vs. 30s)
ax = sns.countplot(data=age_groups, x="Age_Groups", order=['20s', '30s'])

# add labels
for p, label in zip(ax.patches, age_groups['Age_Groups'].value_counts()):
    ax.annotate(label, (p.get_x()+0.37, p.get_height()+0.15))
plt.title("Age Groups (20s vs. 30s)")
plt.ylabel('Number of Customers')
plt.show()
```



Summary

- Before visualizing the target variable overlaid, we can see here that after binning age into two groups, there are more customers who are in their 20s than customers in their 30s in this data set.

4.2.5 Percentage of Travel Insurance Purchases on Various Features

Python code:

```
[10]: def make_stacked_barcharts(df, x):  
    '''  
    Takes in a data frame 'df' and a column 'x'  
    and returns a stacked bar chart of the column  
    with the percentage of purchases overlaid  
    '''  
  
    # Calculate total counts from both groups  
    total = df.groupby(x)['TravelInsurance'].count().reset_index()
```

```

# Calculate total counts from only purchases
purchase = df[df.TravelInsurance == 1].groupby(
    x)['TravelInsurance'].count().reset_index()

# get percentages for purchases
purchase['TravelInsurance'] = [
    i / j * 100
    for i, j in zip(purchase['TravelInsurance'], total['TravelInsurance'])
]

# get percentages
total['TravelInsurance'] = [
    i / j * 100
    for i, j in zip(total['TravelInsurance'], total['TravelInsurance'])
]

# bar chart 1 -> top bars (group of 'TravelInsurance=0')
bar1 = sns.barplot(x, y="TravelInsurance", data=total, color='darkblue')

# bar chart 2 -> bottom bars (group of 'TravelInsurance=1')
bar2 = sns.barplot(x,
    y="TravelInsurance",
    data=purchase,
    color='lightblue')

# add legend
top_bar = mpatches.Patch(color='darkblue', label='Non-Purchase')
bottom_bar = mpatches.Patch(color='lightblue',
    label='Travel Insurance Purchased')
plt.legend(handles=[top_bar, bottom_bar],
    bbox_to_anchor=(1.05, 1),
    loc=2,
    borderaxespad=0.)

# Aesthetics
plt.title("Percentage of Travel Insurance Purchases by " + x)
plt.xlabel(x)
plt.ylabel("% of Purchases")

# Change ticks on x-axis for ChronicDiseases column
if x == "ChronicDiseases":
    chronicdisease = [0, 1]
    labels = ['No', 'Yes']
    plt.xticks(chronicdisease, labels)


# show the graph
plt.show()


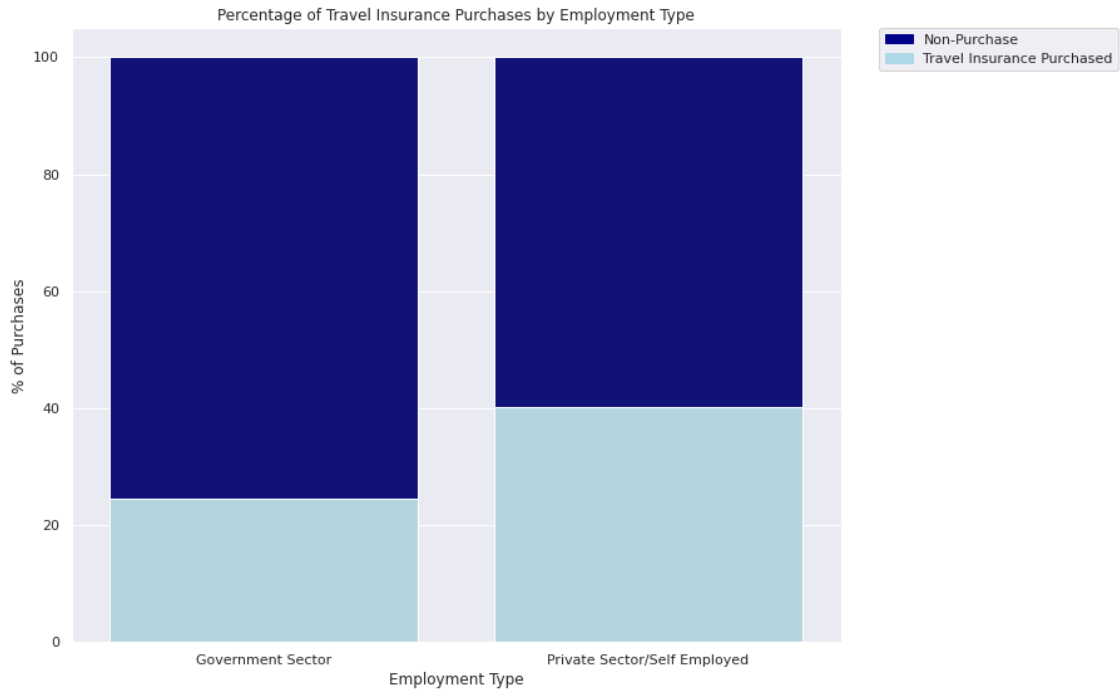
```



```
[11]: # Call stacked bar chart function
make_stacked_barcharts(age_groups, 'Age_Groups')

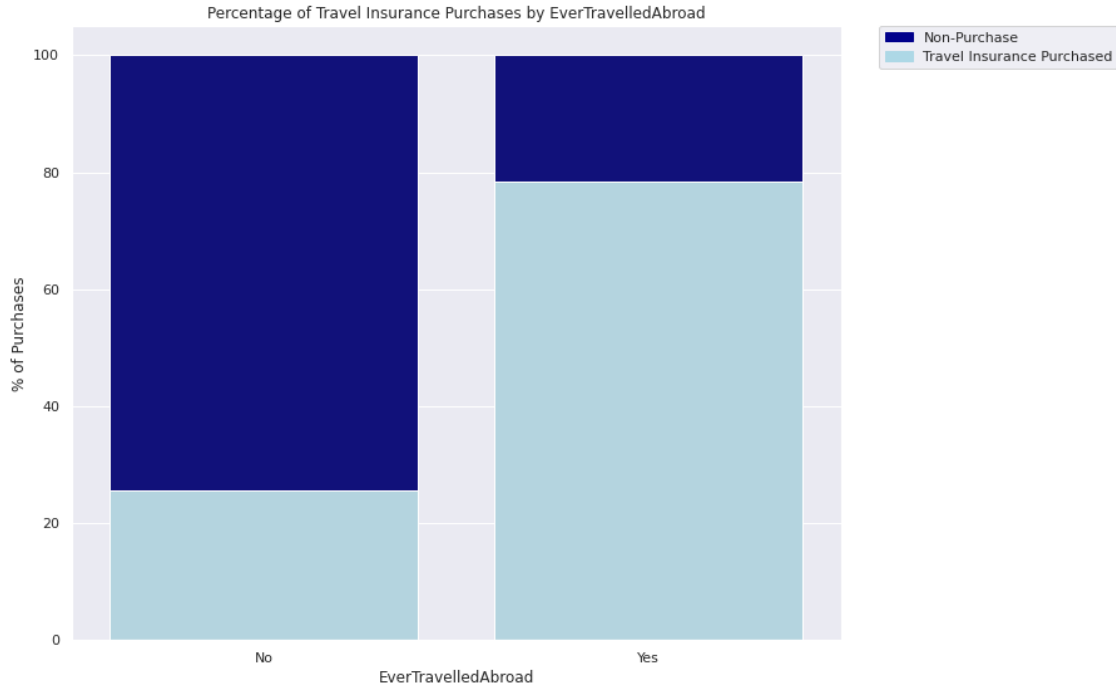
# Make stacked bar charts on various columns
plot_columns = [
    'Employment Type', 'GraduateOrNot', 'ChronicDiseases', 'FrequentFlyer',
    'EverTravelledAbroad'
]

# Plot each column
for i in plot_columns:
    make_stacked_barcharts(df, i)
```









Summary

- There is a higher proportion of customers in their 30s that purchased travel insurance.
- There is a higher proportion of customers who works in a private sector or is self-employed that purchased travel insurance.
- There is no significant difference in proportion between customers who are a college graduate or not that purchased travel insurance.
- This also applies to customers with or without chronic diseases that purchased travel insurance.
- However, there is a higher proportion of customers who are frequent flyers and/or have traveled abroad that purchased travel insurance.

4.3 Monthly Income (Indian Rupee) Range

```
[12]: # Counts by Monthly Rupee
n = np.array([[20., 63., 60., 67., 41., 67., 74., 191., 84., 43.],
              [208., 239., 223., 273., 153., 268., 274., 204., 87., 58.]])

# Bins for Montly Rupee
bins= np.array([ 25000., 37500., 50000., 62500., 75000., 87500., 100000.,
                 112500., 125000., 137500., 150000.])

# Create normalized histogram for groups by target overlay
n_table = np.column_stack((n[0], n[1])) # stack the tables
n_norm = n_table / n_table.sum(
```

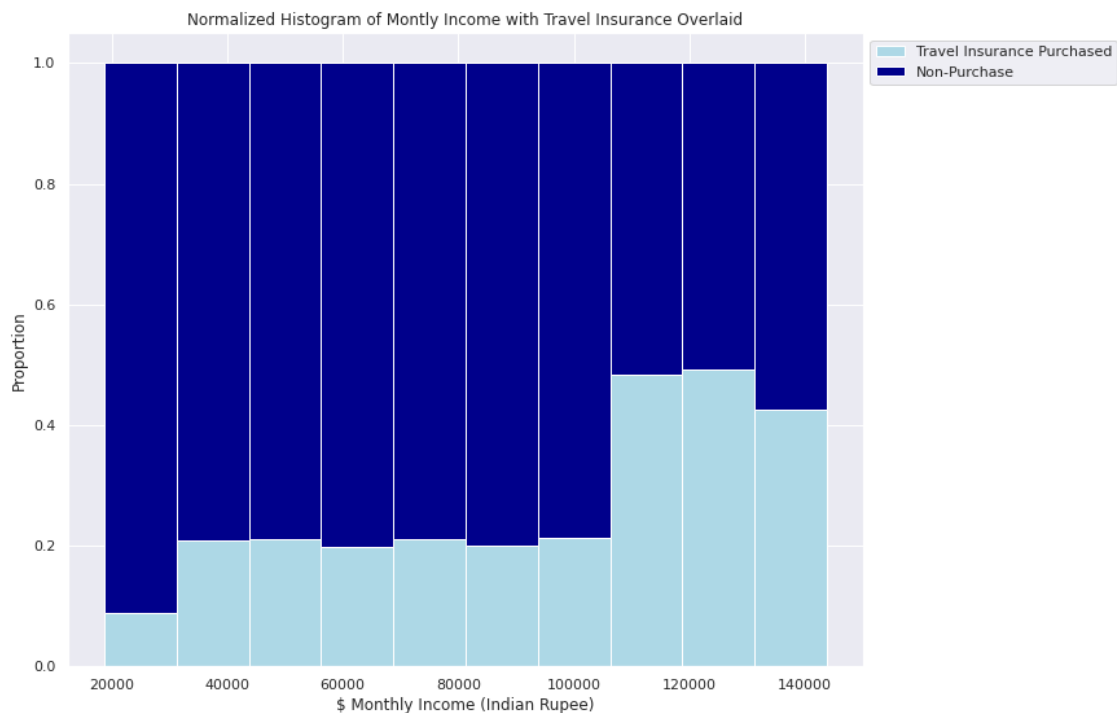


```

axis=1)[: , None] # create normalized tables by sum
ourbins = np.column_stack((bins[0:10], bins[1:11])) # create table bins

p1 = plt.bar(x=ourbins[:, 0],
             height=n_norm[:, 0],
             width=ourbins[:, 1] - ourbins[:, 0],
             color="lightblue") # first bar chart
p2 = plt.bar(
    x=ourbins[:, 0],
    height=n_norm[:, 1],
    width=ourbins[:, 1] - ourbins[:, 0], # second bar chart
    bottom=n_norm[:, 0],
    color="darkblue")
# Annotate legend, title with x and y labels
plt.legend(['Travel Insurance Purchased', 'Non-Purchase'],
           bbox_to_anchor=(1, 1))
plt.title('Normalized Histogram of Montly Income with Travel Insurance_
↳Overlaid')
plt.xlabel('$ Monthly Income (Indian Rupee)')
plt.ylabel('Proportion')
plt.show()

```



4.4 Side-by-side Box-plots between Annual Income and Different Attributes

Python code:

```
[13]: def make_boxplots(df, x):  
    '''  
    Takes in 'x' as a column from data frame 'df'  
    and returns a side by side box-plot of  
    x on the x-axis and AnnualIncome on the y-axis  
    seperated by different colors noted by TravelInsurance  
    '''  
  
    # Palatte to color the target variable  
    palette = {0: "darkblue", 1: "lightblue"}  
  
    # Change x-axis labels if age_groups or GraduatedOrNot  
    order = None  
    if x == "Age_Groups":  
        order = ["20s", "30s"]  
    if x == "GraduateOrNot":  
        order = ["No", "Yes"]  
  
    #Convert AnnualIncome to Monthly  
    df['Monthly_Income'] = round(df['AnnualIncome']/12,2)  
  
    # Boxplot  
    sns.boxplot(x=x,  
                y="Monthly_Income",  
                hue="TravelInsurance",  
                data=df,  
                order=order,  
                palette=palette)  
  
    # Legend properties  
    top_bar = mpatches.Patch(color='darkblue', label='Non-purchase')  
    bottom_bar = mpatches.Patch(color='lightblue',  
                                label='Travel Insurance Purchased')  
    plt.legend(handles=[top_bar, bottom_bar],  
               bbox_to_anchor=(1.05, 1),  
               loc=2,  
               borderaxespad=0.)  
  
    # Graph Properties  
    plt.title(x + " vs. Montly Income with Travel Insurance Overlaid ")  
    plt.xlabel(x)  
    plt.ylabel("$ Monthly Income (Indian Rupee)")
```

```

# Change ticks on x-axis for ChronicDiseases column
if x == "ChronicDiseases":
    chronicdisease = [0, 1]
    labels = ['No', 'Yes']
    plt.xticks(chronicdisease, labels)

# show the graph
plt.show()

```

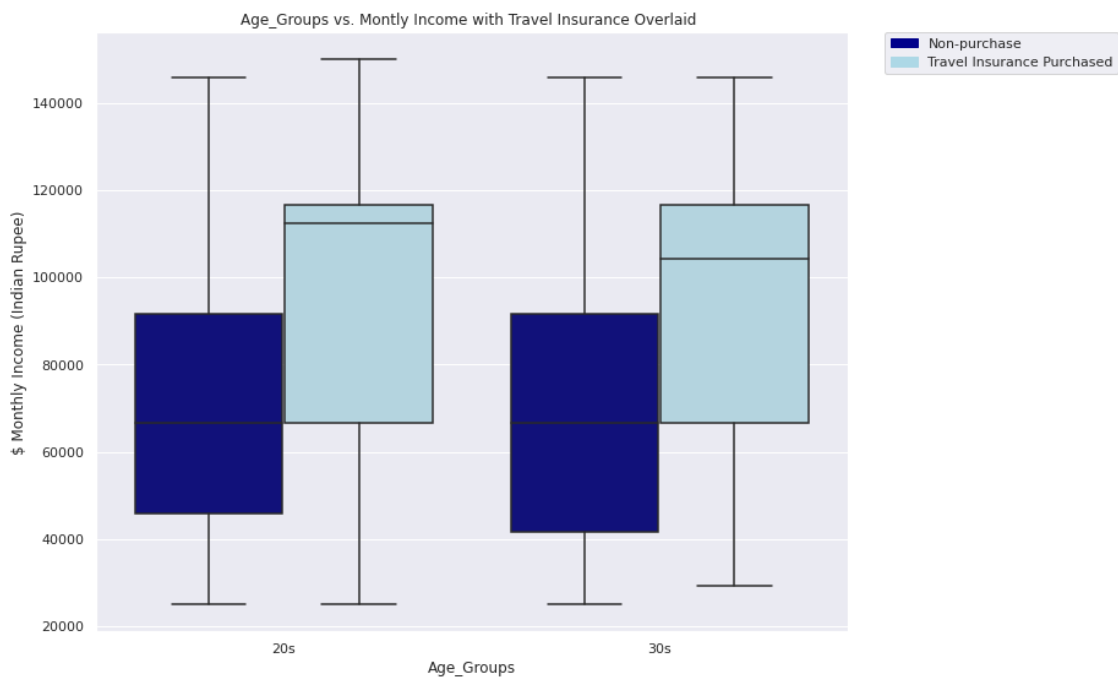
```

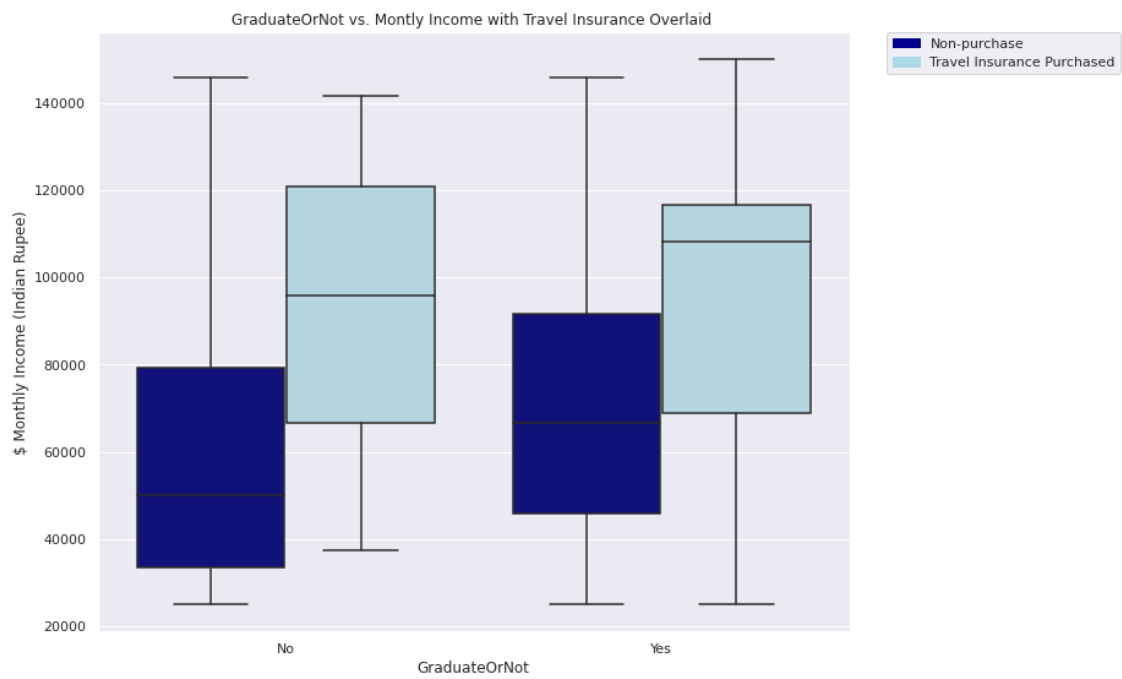

[14]: # call stacked bar chart function
make_boxplots(age_groups, 'Age_Groups')


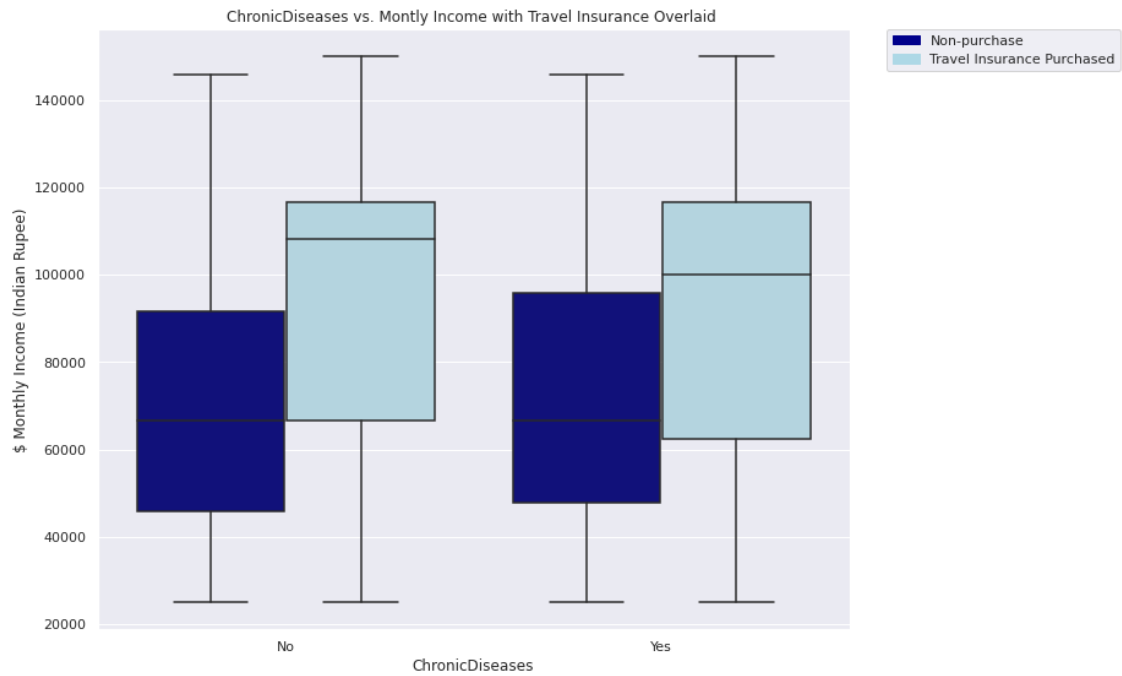
# call boxplot function on various columns
plot_columns = ['Employment Type', 'GraduateOrNot', 'ChronicDiseases',
↳ 'FrequentFlyer',
    'EverTravelledAbroad']


# plot each column
for i in plot_columns:
    # boxplot function
    make_boxplots(df, i)

```









5 Feature Engineering and Pre-Processing

- Invent new columns
- Handle missing values, outliers, correlated features, etc.

Python code:

- Bin age into age groups

```
[15]: # Create function to categorize age groups
def age_groups(x):
    """
    x: This is a value from df['Age']
    returns each as a new categorical value of 20s or 30s
    """
    if x < 30:
        return '20s'
    else:
        return '30s'

# group age into 20s and 30s
df['Age'] = df['Age'].apply(lambda x: age_groups(x))
df = df.rename(columns = {'Age': 'Age_Groups'})
df.head()
```

```
[15]: Age_Groups      Employment Type GraduateOrNot AnnualIncome \
0      30s      Government Sector      Yes      400000
1      30s Private Sector/Self Employed      Yes      1250000
2      30s Private Sector/Self Employed      Yes      500000
3      20s Private Sector/Self Employed      Yes      700000
4      20s Private Sector/Self Employed      Yes      700000

      FamilyMembers ChronicDiseases FrequentFlyer EverTravelledAbroad \
0      6      1      No      No
1      7      0      No      No
2      4      1      No      No
3      3      1      No      No
4      8      1      Yes     No

      TravelInsurance Monthly_Income
0      0      33333.33
1      0      104166.67
2      1      41666.67
3      0      58333.33
4      0      58333.33
```

- Create poor- lower - middle - high income classes

```
[16]: # Create function to categorize wealth groups
def wealth_groups(x):
    '''
    x: This is a value from df['Monthly_Income']
    returns each as a new categorical value of 20s or 30s

    The range for categorizing income class

    - Poor Class: 2500- 6500 per month

    - Lower Class: 6500- 15000 per month

    - Middle Class: 15000- 100000 per month

    - Upper Class: 100000- 350000 per month

    '''

    if x <= 6500:
        return 'poor'
    elif x <= 15000:
        return 'lower'
    elif x <= 100000:
        return 'middle'
```

```

else:
    return 'upper'

# group Monthly Income into wealth groups
df['AnnualIncome'] = df['Monthly_Income'].apply(lambda x: wealth_groups(x))
df = df.drop('Monthly_Income', axis = 1)
df = df.rename(columns = {'AnnualIncome': 'Income_Class'})
df.head()

```

```

[16]: Age_Groups      Employment Type GraduateOrNot Income_Class \
0      30s      Government Sector      Yes      middle
1      30s  Private Sector/Self Employed      Yes      upper
2      30s  Private Sector/Self Employed      Yes      middle
3      20s  Private Sector/Self Employed      Yes      middle
4      20s  Private Sector/Self Employed      Yes      middle

      FamilyMembers  ChronicDiseases FrequentFlyer EverTravelledAbroad \
0                6                1            No            No
1                7                0            No            No
2                4                1            No            No
3                3                1            No            No
4                8                1            Yes            No

      TravelInsurance
0                0
1                0
2                1
3                0
4                0

```

- Convert Family Members to household size categories

```

[17]: # Create function to categorize age groups
def household_groups(x):
    '''
    x: This is a value from df['FamilyMembers']
    returns each as a new categorical value of 20s or 30s

    The range for categorizing income class

    - 1 = single

    - 2-4 = small

    - 5-10 = medium
    '''

```



```

- >10 = large

'''

if x == 1:
    return 'single'
elif x <= 4:
    return 'small'
elif x <= 10:
    return 'medium'
else:
    return 'large'

# group Family Members into household groups
df['FamilyMembers'] = df['FamilyMembers'].apply(lambda x: household_groups(x))
df = df.rename(columns = {'FamilyMembers': 'Household_Size'})
df.head()

```

```

[17]:
Age_Groups      Employment Type GraduateOrNot Income_Class \
0      30s      Government Sector      Yes      middle
1      30s  Private Sector/Self Employed      Yes      upper
2      30s  Private Sector/Self Employed      Yes      middle
3      20s  Private Sector/Self Employed      Yes      middle
4      20s  Private Sector/Self Employed      Yes      middle

Household_Size  ChronicDiseases FrequentFlyer EverTravelledAbroad \
0      medium      1      No      No
1      medium      0      No      No
2      small      1      No      No
3      small      1      No      No
4      medium      1      Yes     No

TravelInsurance
0      0
1      0
2      1
3      0
4      0

```

- Convert Binary Categorical Variable values to 0/1

```

[18]: # Convert Frequent Flyer to 0/1 binary values
df['FrequentFlyer'] = np.where((df['FrequentFlyer'] == 'No'), 0, 1)

```

```

# Convert Ever Traveled Abroad to 0/1 binary values
df['EverTravelledAbroad'] = np.where((df['EverTravelledAbroad'] == 'No'), 0, 1)

# Convert GraduateOrNot to 0/1 binary values
df['GraduateOrNot'] = np.where((df['GraduateOrNot'] == 'No'), 0, 1)

# first few rows
df.head()

```

```

[18]:   Age_Groups      Employment Type  GraduateOrNot  Income_Class \
0      30s      Government Sector              1      middle
1      30s  Private Sector/Self Employed          1      upper
2      30s  Private Sector/Self Employed          1      middle
3      20s  Private Sector/Self Employed          1      middle
4      20s  Private Sector/Self Employed          1      middle

      Household_Size  ChronicDiseases  FrequentFlyer  EverTravelledAbroad \
0      medium              1              0              0
1      medium              0              0              0
2      small              1              0              0
3      small              1              0              0
4      medium              1              1              0

      TravelInsurance
0              0
1              0
2              1
3              0
4              0

```

- One hot encode categorical variables to dummy variables

```

[19]: # one-hot encoding the categorical variables
df = pd.get_dummies(df)
df.head()

```

```

[19]:   GraduateOrNot  ChronicDiseases  FrequentFlyer  EverTravelledAbroad \
0              1              1              0              0
1              1              0              0              0
2              1              1              0              0
3              1              1              0              0
4              1              1              1              0

      TravelInsurance  Age_Groups_20s  Age_Groups_30s \
0              0              0              1
1              0              0              1

```

2	1	0	1
3	0	1	0
4	0	1	0

	Employment Type_Government Sector \
0	1
1	0
2	0
3	0
4	0

	Employment Type_Private Sector/Self Employed	Income_Class_middle \
0	0	1
1	1	0
2	1	1
3	1	1
4	1	1

	Income_Class_upper	Household_Size_medium	Household_Size_small
0	0	1	0
1	1	1	0
2	0	0	1
3	0	0	1
4	0	1	0

6 Data splitting

- Training, validation, and test sets
- Since there does not exist a class imbalance problem, we split the data set into 75% training and 25% validation.

Python code:

```
[20]: # Response Variable
outcome = 'TravelInsurance'
y = df[outcome]

# features - Do not use Target_B or Target_D
predictors = [
    'GraduateOrNot', 'ChronicDiseases', 'FrequentFlyer', 'EverTravelledAbroad',
    'Age_Groups_20s', 'Age_Groups_30s', 'Employment Type_Government Sector',
    'Employment Type_Private Sector/Self Employed', 'Income_Class_middle',
    'Income_Class_upper', 'Household_Size_medium', 'Household_Size_small'
]
X = df[predictors]
```

```
# Set seed to 1 and split on 30% validation
train_X, valid_X, train_y, valid_y = train_test_split(X,
                                                    y,
                                                    test_size=0.25,
                                                    random_state=1)

# Check dimensions
train_X.shape, valid_X.shape
```

[20]: ((1490, 12), (497, 12))

7 Model building strategies

- Describing main research questions and appropriate analytics methods

Python code:

[]:

8 Model performance and hyper-parameter tuning

- Model tuning, comparison, and evaluations

8.1 Decision Tree

Python code:

```
[21]: # user grid search to find optimized tree
param_grid = {
    'max_depth': [5, 10, 15, 20, 25],
    'min_impurity_decrease': [0, 0.001, 0.005, 0.01],
    'min_samples_split': [10, 20, 30, 40, 50],
}

# Run Exhaustive Search
gridSearch = GridSearchCV(DecisionTreeClassifier(random_state=1), param_grid,
    ↪cv=5, n_jobs=-1)
gridSearch.fit(X=train_X, y=train_y)

# Initial Parameters
print('Initial Score: ', gridSearch.best_score_)
print('Initial parameters: ', gridSearch.best_params_)
```

```

# Improving the parameters
param_grid = {
    'max_depth': [3, 4, 5, 6, 7, 8],
    'min_impurity_decrease': [0, 0.001, 0.002, 0.003, 0.005, 0.006, 0.007, 0.
↪008],
    'min_samples_split': [6,7,8,9,10,11,12]
}

# Run Exhaustive Search with fine-tuned parameters
gridSearch = GridSearchCV(DecisionTreeClassifier(random_state=1), param_grid,
↪cv=5, n_jobs=-1)
gridSearch.fit(train_X, train_y)

# Final parameters
print('Improved Score: ', gridSearch.best_score_)
print('Improved parameters: ', gridSearch.best_params_)

# Final Decision Tree
tree_model = gridSearch.best_estimator_

# Fit to Training Data
tree_model.fit(train_X, train_y)

```

Initial Score: 0.7785234899328859

Initial parameters: {'max_depth': 5, 'min_impurity_decrease': 0.005,
'min_samples_split': 10}

Improved Score: 0.7785234899328859

Improved parameters: {'max_depth': 3, 'min_impurity_decrease': 0,
'min_samples_split': 6}

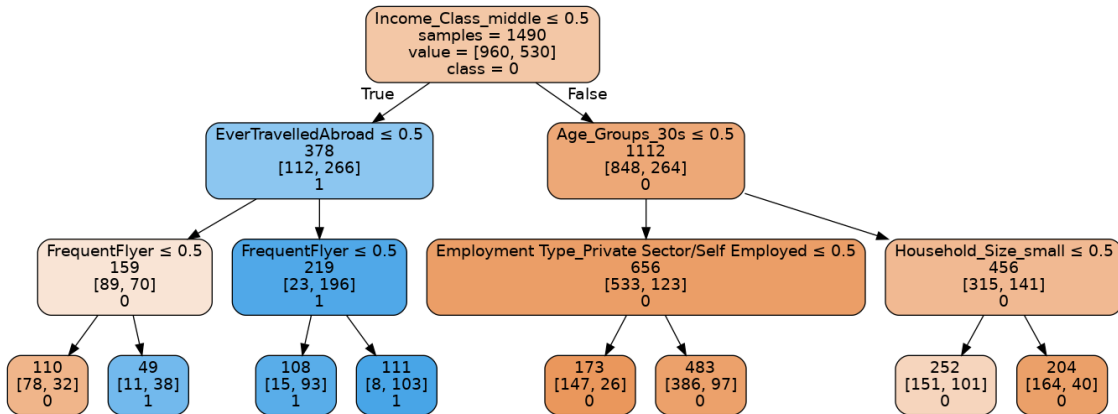
[21]: DecisionTreeClassifier(max_depth=3, min_impurity_decrease=0,
min_samples_split=6, random_state=1)

```

[22]: # Plot Decision tree
plotDecisionTree(tree_model,
                 feature_names=train_X.columns,
                 class_names=tree_model.classes_)

```

[22]:



8.2 Logistic Regression

```
[23]: param_grid = {
    'penalty': ['l1', 'l2'],
    'C': np.logspace(-4, 4, 20),
    'solver': ['liblinear', 'saga']
}

# Create grid search object
gridSearch = GridSearchCV(LogisticRegression(random_state=1, max_iter=5000),
    param_grid=param_grid,
    cv=5,
    n_jobs=-1)

# Fit to training data
gridSearch.fit(train_X, train_y)

# Final parameters
print('Final Score: ', gridSearch.best_score_)
print('Final parameters: ', gridSearch.best_params_)

# Final logistic regression model
logit_model = gridSearch.best_estimator_
```

Final Score: 0.7657718120805369

Final parameters: {'C': 0.004832930238571752, 'penalty': 'l2', 'solver': 'liblinear'}

```
[24]: # Train logistic regression model to find the best predictors
def train_model(variables):
    if len(variables) == 0:
        return None
```

```

model = LogisticRegressionCV(penalty="l2",
                             solver='liblinear',
                             cv=5,
                             random_state=1,
                             max_iter=5000)
return model.fit(train_X[variables], train_y)

# Return the accuracy score in the validation set over each predictor
def score_model(model, variables):
    if len(variables) == 0:
        return 0
    logit_reg_valid = model.predict(valid_X[variables])
    return -accuracy_score(valid_y,
                           [1 if p > 0.5 else 0 for p in logit_reg_valid])

# Use step-wise regression to select the best subset of features
logit_model, best_variables = stepwise_selection(predictors,
                                                train_model,
                                                score_model,
                                                direction='forward',
                                                verbose=True)

print(best_variables)

# Use the previous columns
columns = best_variables

# Fit to Training Data with previous columns
logit_model.fit(train_X[columns], train_y)

```

Variables: GraduateOrNot, ChronicDiseases, FrequentFlyer, EverTravelledAbroad, Age_Groups_20s, Age_Groups_30s, Employment Type_Government Sector, Employment Type_Private Sector/Self Employed, Income_Class_middle, Income_Class_upper, Household_Size_medium, Household_Size_small

Start: score=0.00, constant

Step: score=-0.77, add EverTravelledAbroad

Step: score=-0.79, add Income_Class_middle

Step: score=-0.80, add FrequentFlyer

Step: score=-0.80, unchanged None

['EverTravelledAbroad', 'Income_Class_middle', 'FrequentFlyer']

[24]: LogisticRegressionCV(cv=5, max_iter=5000, random_state=1, solver='liblinear')

8.3 Multi-Layered Neural Network

```
[25]: # user grid search to find optimized hidden layers
param_grid = {
    'hidden_layer_sizes': [(1), (2), (3), (4), (5)],
}

# Run Exhaustive search for neural networks hyper-parameters
gridSearch = GridSearchCV(MLPClassifier(activation = 'logistic',
                                       solver='lbfgs', random_state=1,
                                       ↪max_iter=5000),
                           param_grid,
                           cv=5,
                           n_jobs=-1,
                           return_train_score=True)

# Fit to training set
gridSearch.fit(train_X, train_y)

# Initial Scores and hyper-parameters
print('Initial score: ', gridSearch.best_score_)
print('Initial parameters: ', gridSearch.best_params_)

# Look at Initial Scores with averages
display=['param_hidden_layer_sizes', 'mean_test_score', 'std_test_score']
pd.DataFrame(gridSearch.cv_results_)[display]
```

Initial score: 0.7718120805369127

Initial parameters: {'hidden_layer_sizes': 1}

```
[25]:
```

	param_hidden_layer_sizes	mean_test_score	std_test_score
0	1	0.771812	0.018009
1	2	0.771812	0.008751
2	3	0.771812	0.020022
3	4	0.755705	0.019954
4	5	0.766443	0.020312

```
[26]: # user grid search to fine-tune hyper-parameters
param_grid = {
    'hidden_layer_sizes': [(1, 2), (1, 3), (1, 4), (1, 5), (1, 6)],
}

# Run Exhaustive search for fine-tuning the hyper-parameters
gridSearch = GridSearchCV(MLPClassifier(activation='logistic',
                                       solver='lbfgs',
                                       random_state=1,
                                       max_iter=5000),
                           param_grid,
```



```

        cv=5,
        n_jobs=-1,
        return_train_score=True)

# Fit to training set
gridSearch.fit(train_X, train_y)

# Improved Scores and hyper-parameters
print('Improved score: ', gridSearch.best_score_)
print('Final parameters: ', gridSearch.best_params_)

# Look at fine-tuned hyper-parameters with averages
display = ['param_hidden_layer_sizes', 'mean_test_score', 'std_test_score']
pd.DataFrame(gridSearch.cv_results_)[display]

```

Improved score: 0.7751677852348993

Final parameters: {'hidden_layer_sizes': (1, 3)}

```

[26]: param_hidden_layer_sizes  mean_test_score  std_test_score
0                (1, 2)           0.768456      0.016576
1                (1, 3)           0.775168      0.012006
2                (1, 4)           0.768456      0.018133
3                (1, 5)           0.770470      0.011547
4                (1, 6)           0.771141      0.012979

```

```

[27]: # Final Network
network_model = gridSearch.best_estimator_

# Fit to training data
network_model.fit(train_X, train_y)

```

```

[27]: MLPClassifier(activation='logistic', hidden_layer_sizes=(1, 3), max_iter=5000,
                  random_state=1, solver='lbfgs')

```

```

[28]: # Helper functions
def confusionMatrices(model, title):
    '''
    Takes in a model and the title to return classification
    summary in accuracy and confusion matrix of the model
    '''

    if model == logit_model:
        print(title + ' - Training results')
        classificationSummary(train_y, model.predict(train_X[columns]))
        print(title + ' - Validation results')
        valid_pred = model.predict(valid_X[columns])
        classificationSummary(valid_y, valid_pred)
    else:

```

```

print(title + ' - Training results')
classificationSummary(train_y, model.predict(train_X))
print(title + ' - Validation results')
valid_pred = model.predict(valid_X)
classificationSummary(valid_y, valid_pred)

# Confusion Matrix - Decision Tree
tree_confusion = confusionMatrices(tree_model, '\n\tDecision Tree')

# Confusion Matrix - Logistic Regression
logit_confusion = confusionMatrices(logit_model, '\n\tLogistic regression')

# Confusion Matrix - Multi-layered Neural Network
network_confusion = confusionMatrices(network_model, '\n\tNeural Network')

```

Decision Tree - Training results
Confusion Matrix (Accuracy 0.7785)

	Prediction	
Actual	0	1
0	926	34
1	296	234

Decision Tree - Validation results
Confusion Matrix (Accuracy 0.8089)

	Prediction	
Actual	0	1
0	310	7
1	88	92

Logistic regression - Training results
Confusion Matrix (Accuracy 0.7772)

	Prediction	
Actual	0	1
0	919	41
1	291	239

Logistic regression - Validation results
Confusion Matrix (Accuracy 0.8048)

	Prediction	
Actual	0	1
0	308	9
1	88	92

Neural Network - Training results
Confusion Matrix (Accuracy 0.7805)

	Prediction	
Actual	0	1
0	928	32
1	295	235

Neural Network - Validation results
Confusion Matrix (Accuracy 0.8068)

	Prediction	
Actual	0	1
0	309	8
1	88	92

9 Results and final model selection

- Performance measures on test Set

Python code:

[]:

10 Discussion and conclusion

- Address the problem statement and suggestions that could go beyond the scope of the course

[]: